



All arms trained finitely and reduced PDE loss, but training convergence did not imply event competence

(a) Logged physics residual



(b) Measured V100 cost



(c) Peak GPU allocation

